

CIS 612 BIG DATA FINAL PROJECT REPORT

Niteesh Nitin Singh – 2886321

Dharmik Rajesh Kurlawala – 2886995

❖ Introduction

The QA system leverages Natural Language Processing (NLP) and Neo4j Graph Database to provide real-time, natural language query handling for movie datasets. The project aims to address movie-related queries, extract insights, and deliver an enhanced user experience.

❖ Platform/System Setup

- **Environment:** Python 3.12.3 with essential libraries such as Flask, pandas, pymongo, Neo4j, spacy, tqdm, and others.
- **Python IDE:** Pycharm Community version 2023.3.3
- **Neo4j Configuration:** Requires Neo4j running locally with the credential's user="neo4j", password="bigdata612".

Made few config changes:

- dbms.memory.heap.initial_size=6G
- dbms.memory.heap.max_size=6G
- dbms.memory.pagecache.size=8G
- dbms.windows_service_name=neo4j-relate-dbms-a4ee4310-4e49-49e0-9082-926af4973254

```
# Neo4j Configuration
NEO4J_URI = "bolt://localhost:7687"
NEO4J_USER = "neo4j"
NEO4J_PASSWORD = "bigdata612"
```

- **MongoDB Configuration:** Local MongoDB instance to store intermediate datasets.

```
# MongoDB Setup
client = MongoClient("mongodb://localhost:27017/")
db = client['MoviesDB']
movies_collection = db['movies']
triples_collection = db['triples']
```

- **CoreNLP Server:** Stanza CoreNLP server for natural language processing tasks.
- Required Libraries: pandas, pymongo, psutil, neo4j, Flask, request, render_template.

❖ Setup Procedures

- Install the required packages: flask pandas pymongo spacy tqdm neo4j stanza
- Set up and run Neo4j locally, ensuring it listens on bolt://localhost:7687.
- Download the Stanza CoreNLP package and specify the path in NLP_triple_extract.py.

```
import stanza
stanza.install_corenlp()
```

```
# CoreNLP Path
CORENLP_PATH = "C:/Users/nitee/stanza_corenlp/*"
```

- Load the spaCy model for English: en_core_web_trf

```
# Load the transformer-based English model
nlp = spacy.load('en_core_web_trf')
```

❖ Code Analysis

Each file serves a specific purpose in the project pipeline:

1. Dataset Cleaning (Dataset_cleaning.py)

- **Purpose:** Cleans and processes raw movie data.
- **Steps:**
 - Handle missing values in columns like budget, revenue, runtime, etc.

```
def handle_missing_values(df):
    df = df[df['title'].notna()]
    df = df[df['title'].str.strip() != ""]
    df['budget'] = df['budget'].mask(df['budget'] == 0, pd.NA)
    df['revenue'] = df['revenue'].mask(df['revenue'] == 0, pd.NA)
    df['runtime'] = df['runtime'].mask(df['runtime'] <= 0, pd.NA)
    textual_columns_to_fill = ['genres', 'director', 'imdb_id', 'spoken_languages',
                              'cast', 'production_companies', 'production_countries', 'writers',
                              'director_of_photography', 'producers', 'music_composer']
    for col in textual_columns_to_fill:
        if col in df.columns:
            df[col] = df[col].fillna("unknown")
    if 'overview' in df.columns:
        df['overview'] = df['overview'].str.strip().replace(["", "n/a", "none"], "unknown")
    return df
```

- Standardize text by converting to lowercase and stripping whitespace.

```
1 usage new *
def standardize_text_columns(df, columns):
    for col in tqdm(columns, desc="Standardizing text columns"):
        df[col] = df[col].astype(str).str.strip().str.lower()
    return df
```

- Normalize multi-value columns (e.g., genres, cast) by splitting strings into lists.

```
1 usage new *
def normalize_columns(df):
    columns_to_split = [
        'genres', 'spoken_languages', 'production_companies', 'production_countries',
        'cast', 'director', 'director_of_photography', 'writers', 'producers', 'music_composer'
    ]
    for col in tqdm(columns_to_split, desc="Normalizing columns"):
        if col in df.columns:
            df[col] = df[col].str.strip().str.split(",")
    return df
```

- Remove duplicate records based on title and release_date.

```
3 usages (2 dynamic) new *
def remove_duplicates(df):
    df = df.drop_duplicates(subset=['title', 'release_date'], keep='first')
    return df
```

```
cleaned_data = (
    new_data
    .pipe(lambda df: tqdm.pandas(desc="Handling missing values") or handle_missing_values(df))
    .pipe(lambda df: tqdm.pandas(desc="Standardizing text") or standardize_text_columns(df, textual_columns))
    .pipe(lambda df: tqdm.pandas(desc="Correcting inconsistent data") or correct_inconsistent_data(df))
    .pipe(lambda df: tqdm.pandas(desc="Normalizing columns") or normalize_columns(df))
    .pipe(lambda df: tqdm.pandas(desc="Removing duplicates") or remove_duplicates(df))
)
```

- Save the cleaned dataset to MongoDB in the MoviesDB.movies collection.

```
batch_size = 1000
for i in tqdm(range(0, len(records), batch_size), desc="Inserting records to MongoDB"):
    batch = records[i:i+batch_size]
    collection.insert_many(batch)
```

- **Intermediate Output:** A clean and structured dataset in MongoDB.

```
1 > import ...
4
5 tqdm.pandas()
6
7 file_path = 'TMDB_all_movies_post_2020(200k).csv'
8 new_data = pd.read_csv(file_path)
9
10 textual_columns = [
11     'title', 'status', 'imdb_id', 'original_language', 'original_title',
12     'overview', 'tagline', 'genres', 'production_companies',
13     'production_countries', 'spoken_languages', 'cast', 'director',
14     'director_of_photography', 'writers', 'producers', 'music_composer'
15 ]
16
17
18 1 usage new *
19 def handle_missing_values(df):
20     df = df[df['title'].notna()]
21     df = df[df['title'].str.strip() != '']
22     df['budget'] = df['budget'].mask(df['budget'] == 0, pd.NA)
23     df['revenue'] = df['revenue'].mask(df['revenue'] == 0, pd.NA)
24     df['runtime'] = df['runtime'].mask(df['runtime'] <= 0, pd.NA)
```

Run D:\test\venv\CIS612_Final_Project\Scripts\python.exe D:\CIS612_Final_Project\Dataset_cleaning.py

Standardizing text columns: 100% | 17/17 [00:00<00:00, 21.48it/s]

Normalizing columns: 100% | 19/19 [00:01<00:00, 7.39it/s]

Inserting records to MongoDB: 100% | 206/206 [00:08<00:00, 25.36it/s]

Data successfully imported to MongoDB.

Process finished with exit code 0

2. Triple Extraction (NLP_triple_extract.py)

- **Purpose:** Extracts subject-predicate-object triples from movie overviews using CoreNLP.

- **Steps:**

- Resolves coreferences in text to improve triple extraction accuracy.

```
def resolve_coreferences(annotated):
    if not annotated.corefChain:
        return annotated.text

    replacements = {}
    for chain in annotated.corefChain:
        representative = chain.mention[0] # Representative mention
        rep_text = " ".join(
            token.word
            for token in annotated.sentence[representative.sentenceIndex].token[
                representative.beginIndex:representative.endIndex
            ]
        )

        for mention in chain.mention[1:]:
            sent_idx = mention.sentenceIndex
            token_range = range(mention.beginIndex, mention.endIndex)
            mention_text = " ".join(
                annotated.sentence[sent_idx].token[i].word for i in token_range
            )
            replacements[mention_text] = rep_text

    resolved_text = annotated.text
    for mention_text, rep_text in replacements.items():
        resolved_text = re.sub(r"\b" + re.escape(mention_text) + r"\b", rep_text, resolved_text)

    return resolved_text
```

- Annotates text using CoreNLP and extracts OpenIE triples.
- Saves triples to MongoDB in the MoviesDB.triples collection.

```
client.start()
all_triples = []
for _, row in tqdm(data.iterrows(), total=len(data)):
    if pd.notna(row.get("overview")):
        extracted_triples = process_overview(row["overview"], row["id"], client)
        all_triples.extend(extracted_triples)

# Save to MongoDB
triples_collection = db["triples"]
triples_collection.drop()
triples_collection.insert_many(all_triples)

print("Triple extraction and storage completed.")
```

- **Issues Encountered:** Port conflicts with CoreNLP, resolved by dynamically assigning a free port.

- **Intermediate Output:** Triples extracted and stored in MongoDB.

The screenshot shows a VS Code editor with a project named 'CIS612_Final_Project'. The file explorer on the left lists various Python scripts. The main editor displays the 'NLP_triple_extract.py' file, which contains a function 'process_overview' that extracts OpenIE triples from movie overviews. The 'Run' panel at the bottom shows the command used to execute the script, and the 'Console' panel shows the output, including a progress bar and a completion message.

```

73
74
75 def process_overview(overview, movie_id, client):
76     triples = []
77     ann = client.annotate(overview)
78     resolved_text = resolve_coreferences(ann)
79     resolved_ann = client.annotate(resolved_text)
80
81     # Extract OpenIE triples
82     for sentence in resolved_ann.sentence:
83         for triple in sentence.openieTriple:
84             triples.append({
85                 "movie_id": movie_id,
86                 "subject": triple.subject,
87                 "predicate": triple.relation,
88                 "object": triple.object,
89             })
90
91     return triples
92
93 process_overview(

```

Run: D:\test\.venv\CIS612_Final_Project\Scripts\python.exe D:\CIS612_Final_Project\NLP_triple_extract.py

100% [Progress Bar] 205337/205337 [11:04:34<00:00, 5.15it/s]

Triple extraction and storage completed.

Process finished with exit code 0

3. Export to CSV for Neo4j (Export_to_csv.py)

- **Purpose:** Process cleaned data to CSV for Neo4j import.
- **Steps:**
 - Fetch movies and relationships from MongoDB.

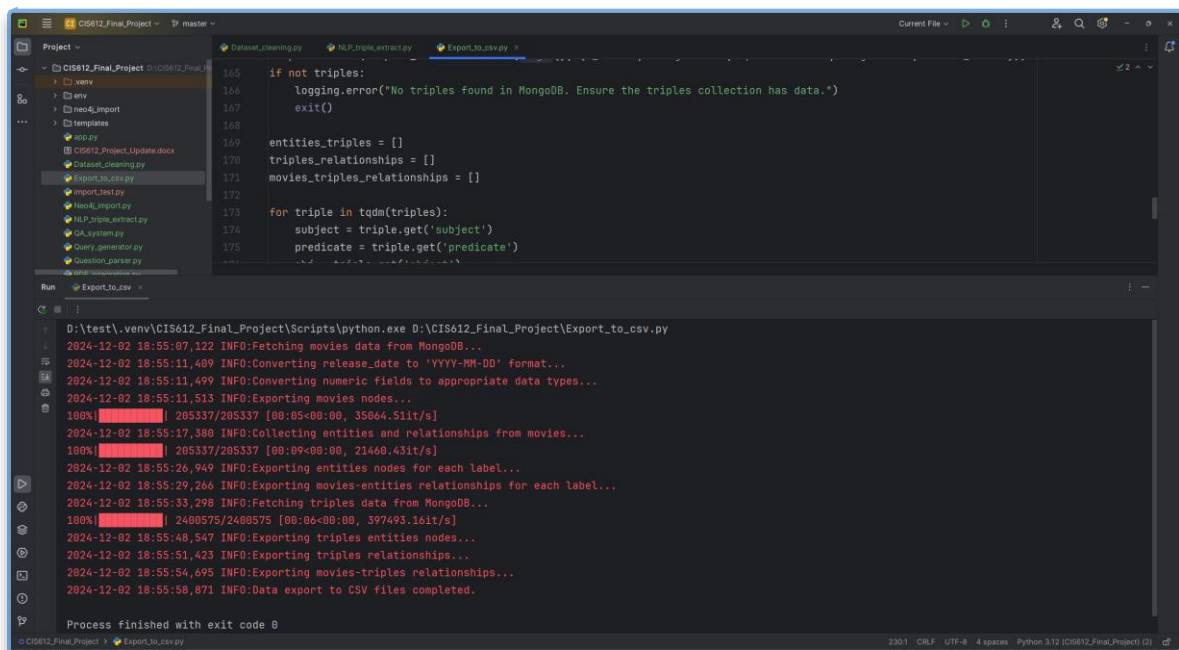
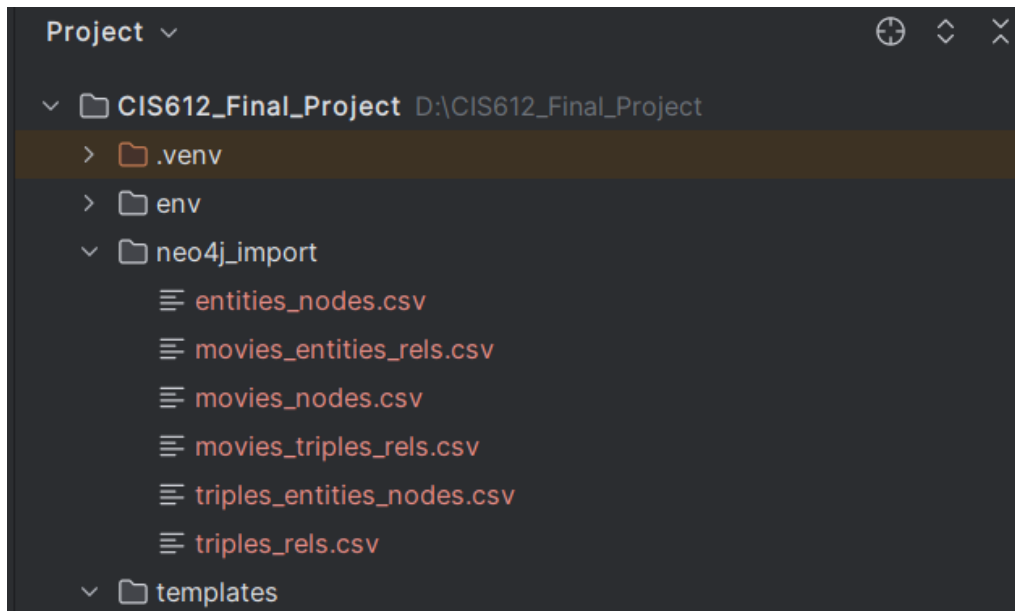
```

# Fetch movies data
logging.info("Fetching movies data from MongoDB...")
movies = list(movies_collection.find())
if not movies:
    logging.error("No movies found in MongoDB. Ensure the movies collection has data.")
    exit()
movies_df = pd.DataFrame(movies)

```

- Normalize names and generate unique entity IDs.
- Create node and relationship files for Neo4j import:
 - Entities nodes
 - Movies entities relations
 - Movies nodes

- Movies triples relations
- Triples entities nodes
- Triple relations
- **Issues Encountered:** Data type mismatches in numeric fields, resolved by coercion and filling NaN values.
- **Intermediate Output:** CSV files stored in the neo4j_import directory.



- **CSV files can be used to perform bulk neo4j-admin import:**

Command:

```
bin\neo4j-admin database import full ^
--nodes=Movie="D:\CIS612_Final_Project\neo4j_import\movies_nodes.csv" ^
--nodes="D:\CIS612_Final_Project\neo4j_import\entities_nodes.csv" ^
--nodes="D:\CIS612_Final_Project\neo4j_import\triples_entities_nodes.csv" ^
--relationships="D:\CIS612_Final_Project\neo4j_import\movies_entities_rels.csv" ^
--relationships="D:\CIS612_Final_Project\neo4j_import\triples_rels.csv" ^
--relationships="D:\CIS612_Final_Project\neo4j_import\movies_triples_rels.csv" ^
--overwrite-destination=true ^
--verbose ^
neo4j
```

Output:

[illegible]


```

Neo4j Desktop Terminal - Graph DBMS
File Edit View Window Help Developer
..... 95X f0827ms [7s 762ms]
..... 100X f084ms [7s 762ms]
Imported 2,472,380 nodes in 7s 70ms
Prepare ID mapper
..... 5X f0897ms [97ms]
..... 10X f085ms [60ms]
..... 15X f086ms [182ms]
..... 20X f086ms [180ms]
..... 25X f086ms [115ms]
..... 30X f087ms [122ms]
..... 35X f0827ms [307ms]
..... 40X f084ms [122ms]
..... 45X f081ms [144ms]
..... 50X f086ms [152ms]
..... 55X f087ms [165ms]
..... 60X f087ms [168ms]
..... 65X f088ms [177ms]
..... 70X f088ms [185ms]
..... 75X f089ms [194ms]
..... 80X f0827ms [172ms]
..... 85X f081ms [153ms]
..... 90X f081ms [164ms]
..... 95X f089ms [157ms]
..... 100X f088ms [165ms]
Range partitioned on node ID: nodeIDrange[numNodes:1, numRangesByNodeData:1, numRangesByRelationshipData:1]
Converting to intermediary format
..... 5X f0817ms [315ms]
..... 10X f0815ms [459ms]
..... 15X f0827ms [137ms]
..... 20X f0809ms [195ms]
..... 25X f0827ms [167ms]
..... 30X f0849ms [197ms]
..... 35X f0809ms [177ms]
..... 40X f087ms [187ms]
..... 45X f088ms [192ms]
..... 50X f0817ms [1s 40ms]
..... 55X f087ms [1s 122ms]
..... 60X f087ms [1s 183ms]
..... 65X f087ms [1s 235ms]
..... 70X f086ms [1s 377ms]
..... 75X f0809ms [1s 383ms]
..... 80X f088ms [1s 48ms]
..... 85X f0817ms [1s 646ms]
..... 90X f0809ms [1s 73ms]
..... 95X f087ms [1s 878ms]
..... 100X f088ms [1s 89ms]
Using Configuration Configuration numberWorkers=20, temporaryPath=C:\Users\ntice\Neo4jDesktop\relate-data\dmms-9c5c83a-11f7-473f-b456-3f3a7ea0c90\data\database\neo4j\temp, numberofTrackedDense=10000, applyBatchSize=64
Switching to bigger page cache for relationship memory 5.2mb
Importing relationships
..... 5X f08957ms [957ms]
..... 10X f0825ms [1s 315ms]
..... 15X f08387ms [1s 621ms]
..... 20X f0825ms [1s 921ms]
..... 25X f08229ms [1s 158ms]
..... 30X f0827ms [1s 433ms]
..... 35X f0867ms [1s 980ms]
..... 40X f0838ms [1s 540ms]
..... 45X f0850ms [1s 527ms]
..... 50X f0873ms [1s 257ms]
..... 55X f0857ms [1s 83ms]
..... 60X f0858ms [1s 484ms]
..... 65X f0857ms [1s 804ms]
..... 70X f0867ms [1s 497ms]
..... 75X f0817ms [1s 85ms]
..... 80X f083ms [1s 165ms]
..... 85X f0817ms [1s 75ms]
..... 90X f083ms [1s 187ms]
..... 95X f0837ms [1s 428ms]
..... 100X f0827ms [1s 721ms]
Imported 2,472,380 relationships in 11s 90ms
Flushing stores
.....

```

4. Neo4j Import (Neo4j_import.py -for small data set)

- **Purpose:** Imports nodes and relationships from MongoDB into Neo4j.
- **Steps:**
 - Creates movie nodes with properties like title, budget, and runtime.

```

def import_movies_and_relationships(movies, neo4j_manager):
    print("Importing movies and relationships into Neo4j...")

    for movie in movies:
        # Movie Node
        movie_id = movie.get("id")
        title = movie.get("title")
        release_date = movie.get("release_date")
        budget = movie.get("budget")
        runtime = movie.get("runtime")
        vote_average = movie.get("vote_average")
        status = movie.get("status")
        revenue = movie.get("revenue")
        original_title = movie.get("original_title")
        imdb_rating = movie.get("imdb_rating")
        imdb_votes = movie.get("imdb_votes")
        original_language = movie.get("original_language")
        popularity = movie.get("popularity")

```

- Establishes relationships such as DIRECTED_BY, BELONGS_TO_GENRE, etc.

```
# Process Relationships
relationships = [
    ("genres", "Genre", "BELONGS_TO_GENRE"),
    ("spoken_languages", "Language", "SPOKEN_IN"),
    ("production_companies", "Company", "PRODUCTION_COMPANY"),
    ("production_countries", "Country", "PRODUCED_IN"),
    ("cast", "Person", "ACTED_IN"),
    ("director", "Person", "DIRECTED_BY"),
    ("writers", "Person", "WRITTEN_BY"),
    ("producers", "Person", "PRODUCED_BY"),
    ("music_composer", "Person", "COMPOSED_BY"),
    ("director_of_photography", "Person", "DOP_BY")
]

for field, label, rel_type in relationships:
    entities = movie.get(field, [])
    if not isinstance(entities, list):
        entities = [entities]

    for entity in entities:
        if not entity:
            continue

        query_entity = f"""
        MERGE (n:{label} {{name: $entity_name}})
        MERGE (m:Movie {{id: $movie_id}})
        MERGE (m)-[:{rel_type}]->(n)
        """
        neo4j_manager.execute_query(query_entity, parameters={
            "entity_name": entity,
            "movie_id": movie_id
        })
```

- Imports extracted triples using APOC procedures.

```
# Import Triples into Neo4j using APOC
1 usage new *
def import_triples(triples, neo4j_manager):
    print("Importing triples into Neo4j...")
    query = """
    CALL apoc.periodic.iterate(
        'UNWIND $triples AS triple RETURN triple',
        'MERGE (s:Entity {name: triple.subject})
        MERGE (o:Entity {name: triple.object})
        MERGE (m:Movie {id: triple.movie_id})
        MERGE (s)-[r:RELATION {type: triple.predicate}]->(o)
        MERGE (m)-[:HAS_SUBJECT]->(s)
        MERGE (m)-[:HAS_OBJECT]->(o)',
        {batchSize: 1000, parallel: true, params: {triples: $triples}}
    )
    """
    neo4j_manager.execute_query(query, parameters={"triples": triples})
    print("Triples successfully imported into Neo4j.")

# Execute the import process
try:
    import_movies_and_relationships(movies, neo4j_manager)
    import_triples(triples, neo4j_manager)
finally:
    neo4j_manager.close()
```

- **Issues Encountered:** Missing data in triples, resolved by skipping incomplete records.

5. Question Parsing (Question_parser.py)

- **Purpose:** Determines user intent and extracts relevant entities from natural language questions.
- **Key Features:**
 - Regex-based patterns to identify intents like FindDirector, FindMoviesByGenre.

```
# Define intent patterns
INTENT_PATTERNS = [
    {
        'intent': 'FindDirector',
        'pattern': r'who\s+(directed|is the director of)\s+(?P<Movie>.+)',
        'entities': ['Movie']
    },
    {
        'intent': 'FindActors',
        'pattern': r'who\s+(acted in|starred in|are the actors in)\s+(?P<Movie>.+)',
        'entities': ['Movie']
    },
    {
        'intent': 'FindMoviesByGenre',
        'pattern': r'(which|what)\s+movies\s+(are|belong to|fall under|classified as)\s+(?P<Genre>.+)\s+genre',
        'entities': ['Genre']
    },
    {
        'intent': 'FindMoviesByDirector',
        'pattern': r'which movies\s+(did|were)\s+(?P<Person>.+)\s+(direct|directed)',
        'entities': ['Person']
    }
]
```

- Fallback to spaCy NER for entity recognition.

```
def parse_question(question):
    question = question.lower().strip('?').strip()
    for pattern_info in INTENT_PATTERNS:
        match = re.match(pattern_info['pattern'], question)
        if match:
            intent = pattern_info['intent']
            entities = {entity: match.group(entity).strip().lower() for entity in pattern_info['entities']}
            return intent, entities

    # Fallback to spaCy NER
    doc = nlp(question)
    entities = {}
    for ent in doc.ents:
        if ent.label_ == 'PERSON':
            entities['Person'] = ent.text
        elif ent.label_ in ['WORK_OF_ART', 'MOVIE']:
            entities['Movie'] = ent.text
        elif ent.label_ == 'LANGUAGE':
            entities['Language'] = ent.text
        elif ent.label_ == 'ORG':
            entities['Company'] = ent.text
        elif ent.label_ == 'GPE':
            entities['Country'] = ent.text
        elif ent.label_ == 'NORP':
            entities['Genre'] = ent.text
    intent = 'Unknown' if not entities else 'FindInformation'
    return intent, entities
```

- **Intermediate Output:** Intent and entities (e.g., {'intent': 'FindDirector', 'entities': {'Movie': 'Inception'}}).

6. Query Generation (Query_generator.py)

- **Purpose:** Generates and executes Cypher queries based on user intent.

- **Features:**

- Predefined query templates for intents like FindActors, FindRevenue.

```
def generate_query(self, intent, entities):
    if intent == 'FindDirector':
        movie_title = entities.get('Movie')
        query = """
        MATCH (m:Movie)<-[:DIRECTED_BY]->(d:Person)
        WHERE toLower(m.title) = toLower($movie_title)
        RETURN d.name AS director
        """
        parameters = {'movie_title': movie_title}
        return query, parameters

    elif intent == 'FindActors':
        movie_title = entities.get('Movie')
        query = """
        MATCH (a:Person)<-[:ACTED_IN]->(m:Movie)
        WHERE toLower(m.title) = toLower($movie_title)
        RETURN a.name AS actor
        """
        parameters = {'movie_title': movie_title}
        return query, parameters
```

- Fetches results from Neo4j and formats responses.

```
def get_response(self, intent, entities):
    query, parameters = self.generate_query(intent, entities)
    if not query:
        return "I'm sorry, I couldn't understand your question."

    results = self.execute_query(query, parameters)
    if not results:
        return "I'm sorry, I couldn't find any results."

    if intent == 'FindDirector':
        directors = [record['director'] for record in results]
        movie_title = entities.get('Movie').title()
        directors_list = ', '.join(director.title() for director in directors)
        return f"The director of '{movie_title}' is {directors_list}."

    elif intent == 'FindActors':
        actors = [record['actor'] for record in results]
        movie_title = entities.get('Movie').title()
        actors_list = ', '.join(actor.title() for actor in actors)
        return f"The actors in '{movie_title}' are: {actors_list}."
```

- **Issues Encountered:** Complex queries for trends required optimization by batching and limiting results.
-

7. QA System (QA_system.py and app.py)

- **Purpose:** Provides an interface for users to interact with the system.
- **Modes:**
 - **Command-line Mode (QA_system.py):** Accepts user questions and prints responses.

```
from Question_parser import parse_question
from Query_generator import QueryGenerator

new *
def main():
    generator = QueryGenerator(password="bigdata612")
    print("Welcome to the Movie QA System! Type 'exit' to quit.\n")

    while True:
        question = input("You: ")
        if question.lower() in ['exit', 'quit']:
            break

        intent, entities = parse_question(question)
        if intent == 'Unknown':
            print("Bot: I'm sorry, I couldn't understand your question.")
            continue

        response = generator.get_response(intent, entities)
        print(f"Bot: {response}")

    generator.close()
    print("Goodbye!")

> if __name__ == '__main__':
    ⚡ main()
```

- **Web Interface (app.py):** Flask application with a form-based UI.

```
app = Flask(__name__)
generator = QueryGenerator(password="bigdata612")

new *
@app.route(rule: '/', methods=['GET', 'POST'])
def home():
    if request.method == 'POST':
        question = request.form['question']
        intent, entities = parse_question(question)
        if intent == 'Unknown':
            response = "I'm sorry, I couldn't understand your question."
        else:
            response = generator.get_response(intent, entities)
        return render_template(template_name_or_list: 'index.html', question=question, response=response)
    return render_template('index.html')

2 usages (2 dynamic) new *
@app.route('/shutdown')
def shutdown():
    generator.close()
    func = request.environ.get('werkzeug.server.shutdown')
    if func is None:
        raise RuntimeError('Not running with the Werkzeug Server')
    func()
    return 'Server shutting down...'

if __name__ == '__main__':
    app.run(debug=True, port=8081)
```

- **Execution:**

The image displays a VS Code editor interface with three main panels. The top panel shows the file explorer with a project structure for 'CIS612_Final_Project'. The middle panel shows the code for 'QA_system.py', which includes a 'main()' function that initializes a 'QueryGenerator' and a 'while True' loop. The bottom panel shows the terminal output of the application, displaying a welcome message and a list of languages with their average ratings. The right panel shows the code for 'app.py', which includes a 'Flask' app, a 'home()' route, and a 'def home()' function that handles POST requests by parsing a question and returning a response.

```

def main():
    generator = QueryGenerator(password="bigdata612")
    print("Welcome to the Movie QA System! Type 'exit' to quit.\n")
    while True:
        main()

```

```

import flask
app = flask.Flask(__name__)
generator = QueryGenerator(password="bigdata612")

@app.route('/', methods=['GET', 'POST'])
def home():
    if request.method == 'POST':
        question = request.form['question']
        intent, entities = parse_question(question)
        if intent == 'Unknown':
            response = "I'm sorry, I couldn't understand your question."

```

Terminal Output:

```

D:\test\.venv\CIS612_Final_Project\Scripts\python.exe D:\CIS612_Final_Project\QA_system.py
Welcome to the Movie QA System! Type 'exit' to quit.

You: who directed gunner?
Bot: The director of 'Gunner' is Dimitri Logothetis.
You: which languages have the highest-rated movies
Bot: Languages with the highest-rated movies:
kirundi (Avg. Rating: 6.33)
éwégbé (Avg. Rating: 4.58)
latin (Avg. Rating: 4.32)
română (Avg. Rating: 3.73)
bokmål (Avg. Rating: 3.64)
esperanto (Avg. Rating: 3.33)
kinyarwanda (Avg. Rating: 3.31)
(Avg. Rating: 3.30) پشتو
dansk (Avg. Rating: 3.28)
fulfulde (Avg. Rating: 3.20)
bamanankan (Avg. Rating: 3.10)
français (Avg. Rating: 3.05)
አማርኛ (Avg. Rating: 3.00)
ᱵᱚᱠᱟᱨ (Avg. Rating: 2.97)
èdè yorùbá (Avg. Rating: 2.91)
italiano (Avg. Rating: 2.86)
gaeilge (Avg. Rating: 2.84)
deutsch (Avg. Rating: 2.73)

```


❖ Problems and Resolutions

- **CoreNLP Port Conflicts:** Resolved by dynamic port allocation.

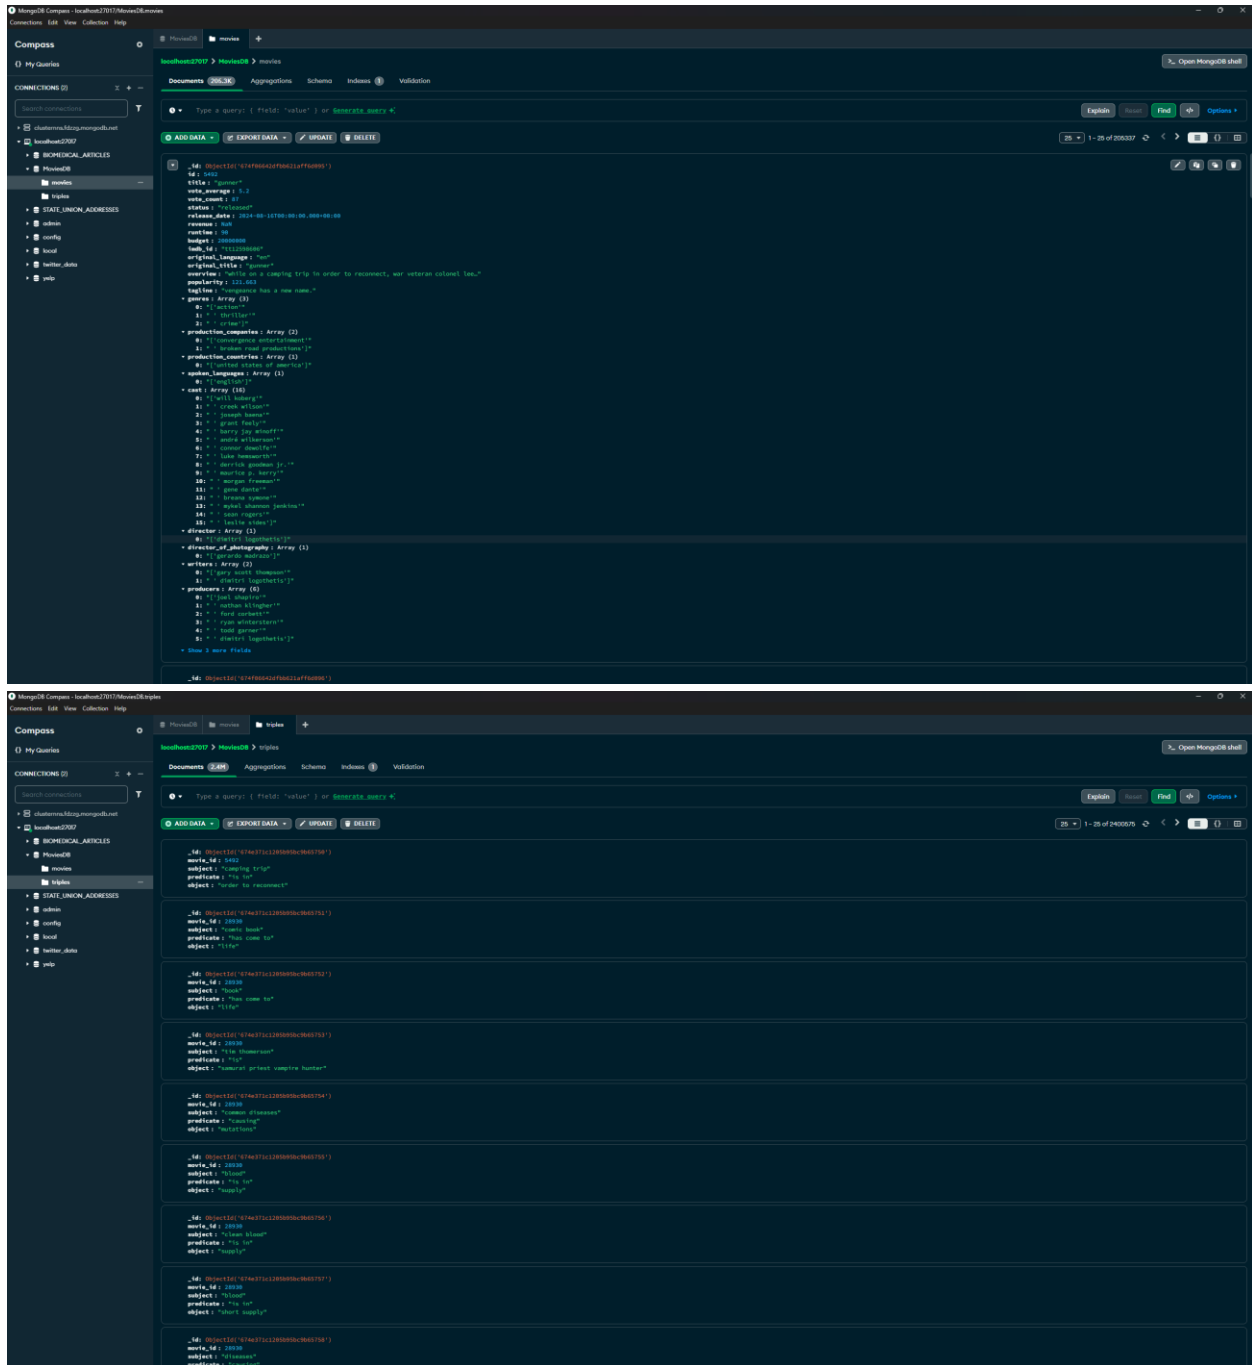
```
def find_free_port():
    with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
        s.bind(('', 0))
        return s.getsockname()[1]

1 usage new *
def kill_corenlp_process(port=9000):
    for proc in psutil.process_iter():
        try:
            for conn in proc.connections(kind="inet"):
                if conn.laddr.port == port:
                    proc.terminate()
        except Exception:
            pass
```

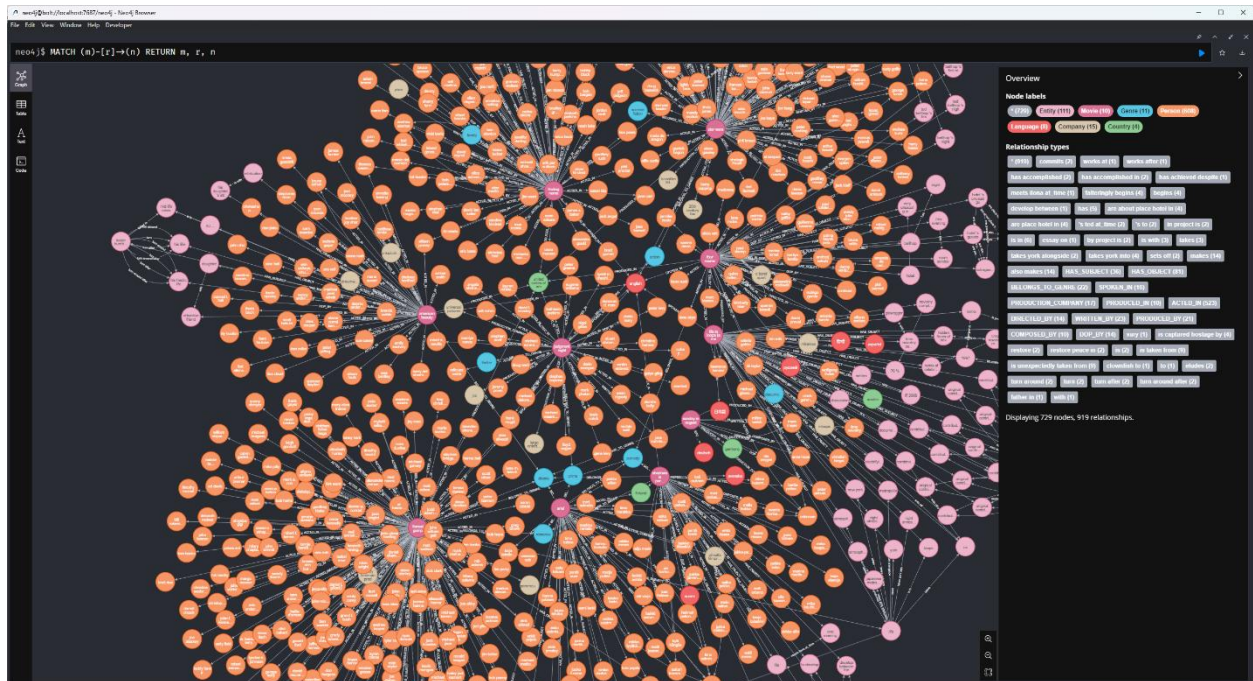
- **Data Type Issues:** Resolved by explicit typecasting in Dataset_cleaning.py and Export_to_csv.py.
- **Neo4j import issue:** Importing a large dataset (~2.2M nodes, ~7.8M relationships, ~211K relationship types) caused bottlenecks.
 - Using Python Driver Import was not feasible due to prolonged import times and performance issues.
 - Solution: Neo4j Admin Import due to its ability to import large datasets in under a minute.
- **Building the QA System issue:** We planned to use Neosemantics (n10s) to build the QA system using SPARQL, which would enable semantic queries on RDF data.
 - The newer version of Neo4j no longer supports Neosemantics, limiting our ability to integrate RDF data and execute SPARQL queries.
 - Solution: Adapted the system to use Cypher Queries, leveraging Neo4j's native query language to design the Q&A functionality.

❖ Final Output

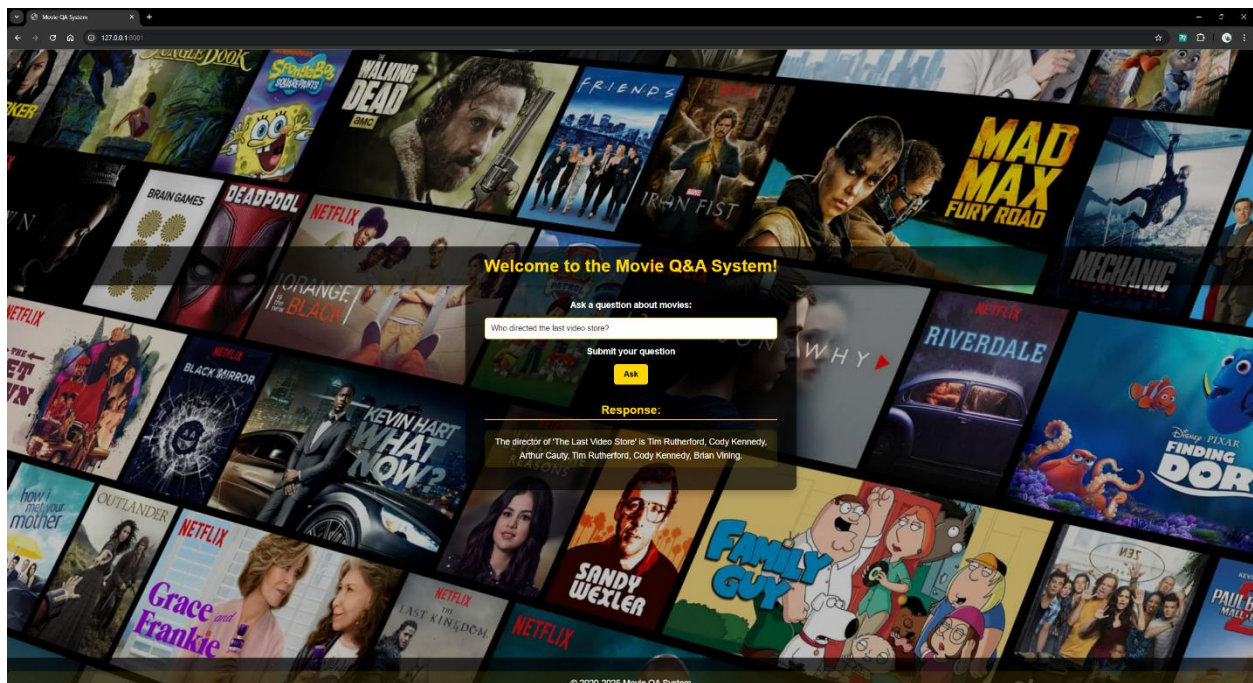
- Cleaned dataset stored in MongoDB.

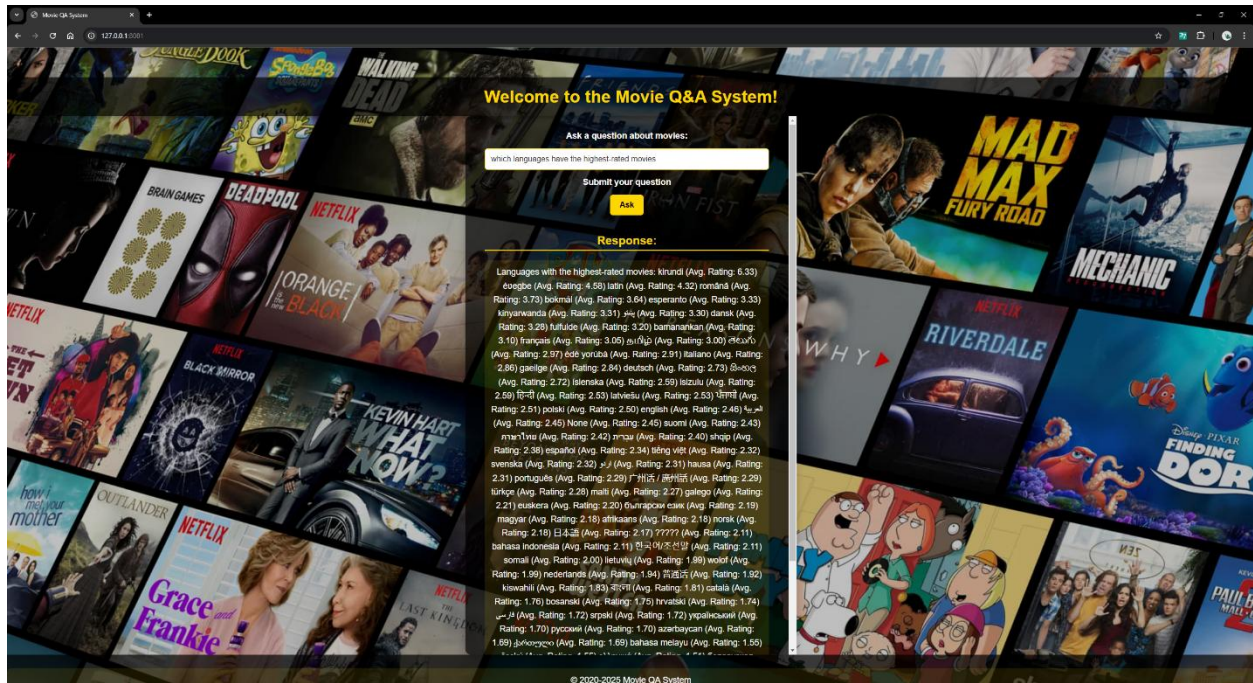


- Knowledge graph nodes and relationships imported into Neo4j.



- A functional QA system capable of answering movie-related questions using the knowledge graph.





❖ Conclusion

- Successfully integrated NLP and graph databases for robust query handling.
- Enhanced natural language understanding with CoreNLP and spaCy.
- Demonstrated scalability for larger datasets and real-world applications.